

Aging, Secular Stagnation and the Business Cycle^{*}

Callum Jones[†]

January 2020

Abstract

As of 2019, US output was 18% below the level predicted by its pre-2007 trend. To understand why, I develop and estimate a model of the US incorporating demographics, real and monetary shocks, and the occasionally binding zero lower bound on nominal rates. Demographic shocks generate slow-moving trends in the real interest rate, employment, and productivity. I find that demographics alone can explain one-third of the gap between log output per capita and its linear trend in 2015. Changing demographics also lowered real rates, causing the zero lower bound to bind between 2009 and 2015, contributing to the slow recovery from the Great Recession.

Keywords: Great Recession, Demographics, Zero Lower Bound, Forward Guidance.

JEL classifications: E3, E4, E5, J11.

1 Introduction

Three significant trends observed over recent decades characterize the US macroeconomic landscape. First, average real growth during expansions has slowed markedly, falling from 3.6% between 1991 and 2001, to 2.9% between 2002 and 2007, and 2.3% since 2009. As a result, by 2019Q3 output remains 18% lower than what is predicted by its long-run pre-Great Recession trend. Second, real and nominal interest rates have fallen, while the Fed Funds rate was at its lower bound between 2009 and 2015. Third, the employment-population ratio has fallen significantly from its peak in the 2000s (Figure 1).

^{*}I am particularly thankful to Virgiliu Midrigan for valuable advice. I thank Saki Bigio, Jaroslav Borovicka, Katarina Borovickova, Tom Cooley, Tim Cogley, Mark Gertler, Dan Greenwald, Mariano Kulish, Julian Kozlowski, Steven Penning, and Thomas Philippon for helpful comments and discussions. I also thank participants of a number of seminars for helpful suggestions. The Appendix is available online at <https://callumjones.github.io>. All remaining errors are mine alone. The views expressed herein are those of the authors and should not be attributed to the IMF, its Executive Board, or its management.

[†]International Monetary Fund, jonescallum@gmail.com

Against this macroeconomic background, the US population has aged. The average age of the population has risen from 28 to 38 from 1970 to 2018. A number of papers have shown that an aging population can explain the trends in growth and interest rates that are observed. [Eggertsson et al. \(2019\)](#) and [Summers \(2014\)](#) show that an older US population can explain why real interest rates are low and therefore why the zero lower bound (ZLB) binds – because savings behavior changes with age, an economy with a higher fraction of older people has more savings and lower interest rates (see also [Carvalho et al., 2016](#); [Gagnon et al., 2016](#)). [Feyrer \(2007\)](#) and [Aaronson et al. \(2014\)](#) show that an aging population can rationalize why productivity growth and the employment-population ratio are low, since younger workers face a steeper human capital profile and older cohorts work fewer hours.

In this paper, I embed demographic shocks into a quantitative New Keynesian model to provide new estimates of the macroeconomic impact of an aging population and to quantitatively explore the interactions between aging, monetary policy, and business cycle shocks. My model is calibrated to household lifecycle data and demographic data, and is estimated using Bayesian methods on aggregate data that include the period of the ZLB. Using the model, I decompose the evolution of aggregate variables into demographic and business cycle shocks.

The existing literature studies the impact of demographic changes by feeding demographic shocks into large-scale overlapping generations models calibrated to match lifecycle and macro moments. My paper instead treats demographic shocks – to fertility and longevity – as one of a set of exogenous drivers of aggregate variables and uses likelihood methods to estimate the model. Under this approach, my model can account for the highly nonlinear interactions between demographic trends, business cycle shocks, and constrained monetary policy. Thus, I provide a full account of the aggregate consequences of an aging population.

I find that demographic shocks alone are responsible for about one-third of the decline in output relative to trend by 2015. I also find quantitatively important nonlinear interactions between demographic changes and business cycle shocks such that, had demographics not changed between 1984 and 2015, the ZLB would not have been a binding constraint on monetary policy. Mitigating the contractionary effect of the ZLB, I find that unconventional monetary policy kept output from falling by a further 2 percentage points between 2011 and mid-2013.

Methodologically, I show how to parsimoniously model anticipated demographic shocks alongside the ZLB and a standard set of business cycle shocks such that it becomes feasible to estimate the model's parameters. In particular, I show that my model is approximated well by an aggregate representation

with time-varying and anticipated changes to the parameters of the model that are exogenous functions of demographic variables. I exploit the resulting computational advantages to filter quarterly output and consumption for the shocks that generate those series, while also accounting for the trends caused by demographic shocks, and for the stimulatory effects of forward guidance by disciplining, with survey data, expectations of how long the ZLB is expected to last each quarter between 2009 and 2015, following [Kulish et al. \(2017\)](#), [Guerrieri and Iacoviello \(2015\)](#), and [Jones \(2017\)](#).

I find that demographic shocks reproduce the long-run trends observed in the US economy, consistent with [Gagnon et al. \(2016\)](#), [Eggertsson et al. \(2019\)](#), and [Aksoy et al. \(2019\)](#). Between 1990 and 2015, changes in the composition of the workforce caused by the aging of the population explain a decline of around 2 percentage points in the labor force participation rate, with a further decline of 4 percentage points expected by 2035. This contraction of the labor force, together with a decline in savings and investment as fewer workers save for retirement, drags down output growth. Furthermore, demographic changes are responsible for about a 1 percentage point decline in real interest rates and about a 2 percentage point decline in the nominal interest rate between 1990 and 2015.

The nonlinearities associated with the ZLB and forward guidance policy mean that it is a nontrivial problem to assess which shocks, other than demographics, were responsible for the decline in output following the Great Recession. To shed light on this, I hold fixed the expected duration of the ZLB each quarter between 2009 and 2015, and find that the contribution of aggregate shocks that capture financial distress depressed output, but were partly offset by positive government spending shocks between 2009 and 2012.

Estimating the model with demographic shocks affects the posterior estimates of the model's structural parameters and the interpretation of which shocks generate business cycles. The estimate of annual trend growth is 0.4% lower when demographics are accounted for, as the changing composition of the workforce implied by those trends is favorable to growth across the estimation sample. Accounting for slow-moving demographic trends also lowers the importance of productivity shocks and raises the role of markup shocks for driving output and consumption. My paper thus relates to studies of how the propagation of shocks change as an economy undergoes structural changes, either anticipated or unanticipated ([Kulish and Pagan, 2016](#); [Canova et al., 2015](#); [Wong, 2015](#); [Jaimovich and Siu, 2012](#); [Fernández-Villaverde et al., 2007](#)). I build on this work by showing how demographic changes give rise to an economy with anticipated time-varying parameters.

2 Model

This section outlines a New Keynesian model with demographic and business cycle shocks. The model features individuals of different ages, monopolistically competitive firms that produce with capital and labor and face price adjustment costs, aggregate shocks, monetary policy with nominal rates subject to the ZLB, and fiscal policy.

2.1 Households

Demographics Individuals are members of overlapping cohorts. Each cohort lives for a maximum of T periods, so that the age range of an individual is 0 to $T - 1$. Cohort s is of mass n_t^s and comprised of a continuum of identical members of age s , measured at the start of period t . I abstract from trend population growth, and normalize the initial population size to 1. The total size of the population at time t is:

$$n_t = \sum_{s=0}^{T-1} n_t^s. \quad (1)$$

Each period, a fraction of each cohort dies with the exogenous age-specific mortality rate γ_t^s :

$$n_{t+1}^{s+1} = (1 - \gamma_t^s) n_t^s. \quad (2)$$

These mortality rates are time-varying; for example, permanent decreases in mortality rates imply increases in longevity. A maximum lifespan of T implies $\gamma_t^{T-1} = 1$.

Household problem An individual of age s has the period utility function $u(c_t^s, \ell_t^s)$ and chooses consumption c_t^s , labor supply ℓ_t^s , capital k_t^s , and one-period risk-free bonds b_t^s , to maximize lifetime utility. The value function of an individual of age s at period t is, in recursive form:

$$V_t^s = \max_{\{c_t^s, \ell_t^s, b_t^s, k_t^s\}} \{ \chi_t u(c_t^s, \ell_t^s) + \beta(1 - \gamma_t^s) \mathbb{E}_t V_{t+1}^{s+1} \}, \quad (3)$$

where the expectation is taken with respect to the aggregate shocks, β is the discount factor common to all individuals, and χ_t is an aggregate autoregressive process subject to iid normal shocks:

$$\ln \chi_t = (1 - \rho_\chi) \ln \chi + \rho_\chi \ln \chi_{t-1} + \sigma_\chi \varepsilon_{\chi,t}. \quad (4)$$

Because of mortality risk, individuals have time-varying and age-specific discount factors $\beta(1 - \gamma_t^s)$.

The unintentional bequests made by individuals of a household who die between periods are aggregated and redistributed to the remaining living households of the same cohort. Individuals have age-specific productivities z^s , receive a transfer from the government ξ_t^s , earn a return R_t on last period's bond holdings b_{t-1}^{s-1} , earn the rental rate on capital r_t on last period's capital holdings k_{t-1}^{s-1} , receive $\frac{1}{n_t^s} \frac{d_t}{p_t}$ dividends from firms, and ψ_t^s for the redistributed unintentional bequest.¹ Households also lend resources $b_{k,t}^s$ to capital-producing firms, who provide capital to goods-producing firms and generate a return $R_{k,t+1}$ for households next period. These features imply that the period budget constraint of the individual of age s is:

$$c_t^s + k_t^s + \frac{b_{k,t}^s}{p_t R_{k,t}} + \frac{b_t^s}{p_t R_t} \leq z^s w_t \ell_t^s (1 - \tau_t) + \frac{b_{t-1}^{s-1}}{p_t} + \frac{b_{k,t-1}^{s-1}}{p_t} + \tau_t^s. \quad (5)$$

where $\tau_t^s = \xi_t^s + \psi_t^s + \frac{1}{n_t^s} \frac{d_t}{p_t} - T_t^g$ collects various transfers and a lump-sum tax T_t^g , and p_t is the price level. In the last period of life, the budget constraint is:

$$c_t^T \leq \frac{b_{k,t-1}^{T-1}}{p_t} + \frac{b_{t-1}^{T-1}}{p_t} + \tau_t^T. \quad (6)$$

By assumption, an individual retires fully from the labor market in her last period of life. Individuals are born with zero wealth, so that $b_{k,t}^0 = 0$ and $b_t^0 = 0$ for all t , and nominal bonds are in net zero supply $b_t = \sum_s n_t^s b_t^s = 0$. Substituting in for the unintentional bequests $\psi_t^s = \frac{n_{t-1}^{s-1}}{n_t^s} \left[\frac{b_{t-1}^{s-1}}{p_t} + \frac{b_{k,t-1}^{s-1}}{p_t} \right]$, and denoting the marginal utility of wealth of an individual of age s in time t by λ_t^s , the optimal choice of risk-free bonds implies the standard Euler equation:

$$\mathbb{E}_t \left[\frac{\lambda_{t+1}^{s+1}}{\lambda_t^s} \right] = \frac{1}{\beta} \mathbb{E}_t \left[\frac{\Pi_{t+1}}{R_t} \right], \quad (7)$$

where $\Pi_t = p_t/p_{t-1}$ is the rate of inflation.

2.2 Firms

There are two types of firms in the economy, intermediate goods producing firms, and capital goods producing firms. The latter are standard, so here the focus will be on the intermediate goods producers.

There are a continuum of intermediate goods firms who hire capital and labor from households

¹The redistribution scheme for unintentional bequests scales the return on savings k_{t-1}^{s-1} and b_{t-1}^{s-1} by $1/(1 - \gamma_{t-1}^{s-1})$.

to supply a substitutable good $y_t(i)$ at price $p_t(i)$ to final goods producers, who in turn use a CES technology with elasticity of substitution ξ_t to aggregate the intermediate goods into a final good which it sells to consumers at the price p_t . Aggregate capital is the sum of each cohort's capital savings: $k_t = \sum_s n_t^s k_t^s$, while aggregate labor hired by the firm is in efficiency units of labor $\ell_t = \sum_s z^s n_t^s \ell_t^s$. I assume that firms use the capital of the deceased in production. The production function of firm i is:

$$y_t(i) = \mu_t^{1-\alpha} (k_{t-1}(i))^\alpha (Z_t \ell_t(i))^{1-\alpha}, \quad (8)$$

where $\frac{Z_t}{Z_{t-1}} = z$ generates trend growth, and μ_t is an aggregate autoregressive TFP process:

$$\ln \mu_t = (1 - \rho_\mu) \ln \mu + \rho_\mu \ln \mu_{t-1} + \sigma_\mu \varepsilon_{\mu,t}, \quad (9)$$

where $\varepsilon_{\mu,t}$ is a standard normal innovation. Denoting $\text{mc}_t(i)$ as the Lagrange multiplier on the firm's cost minimization problem $\min_{\{k_{t-1}(i), \ell_t(i)\}} r_t k_{t-1}(i) + w_t \ell_t(i)$, the rental rate on capital is $r_t = \alpha \text{mc}_t(i) \frac{y_t(i)}{k_{t-1}(i)}$, and the wage is $w_t = (1 - \alpha) \text{mc}_t(i) \frac{y_t(i)}{\ell_t(i)}$. Intermediate goods-producing firms also face a Rotemberg quadratic cost of adjusting prices, parameterized by ϕ_p . The problem of the firm i is to choose its price $p_t(i)$ to maximize firm value:

$$\max_{p_t(i)} \mathbb{E}_t \sum_{t=0}^{\infty} \beta^t \lambda_t \left(\frac{d_t(i)}{p_t} \right), \quad (10)$$

where λ_t is a weighted average of households' shadow values of wealth, and where real dividends are:

$$\frac{d_t(i)}{p_t} = \left(\frac{p_t(i)}{p_t} \right)^{1-\xi_t} y_t - \text{mc}_t(i) \left(\frac{p_t(i)}{p_t} \right)^{-\xi_t} y_t - \frac{\phi_p}{2} \left[\frac{1}{\Pi^*} \frac{p_t(i)}{p_{t-1}(i)} - 1 \right]^2 y_t, \quad (11)$$

where the elasticity of substitution between intermediate goods ξ_t is subject to stochastic shocks, which generates time-varying markups over marginal costs:

$$\ln \xi_t = (1 - \rho_\xi) \ln \xi + \rho_\xi \ln \xi_{t-1} + \sigma_\xi \varepsilon_{\xi,t}. \quad (12)$$

The first order condition on the optimal choice of price resetting is:

$$\beta \phi_p \mathbb{E}_t \frac{\lambda_{t+1}}{\lambda_t} \frac{y_{t+1}}{y_t} \left[\frac{\Pi_{t+1}}{\Pi^*} - 1 \right] \left[\frac{\Pi_{t+1}}{\Pi^*} \right] = \xi_t - 1 - \xi_t \text{mc}_t + \phi_p \left[\frac{\Pi_t}{\Pi^*} - 1 \right] \left[\frac{\Pi_t}{\Pi^*} \right], \quad (13)$$

Log-linearizing (13) gives rise to a forward-looking New Keynesian Phillips curve. I denote the slope of the log-linearized Phillips curve by ϵ_p , which is a function of the steady-state elasticity of substitution ξ and the parameter governing the capital adjustment cost ϕ_p .

The problem of the capital-producing firms is standard: they accumulate capital to maximize market value. Capital adjustment requires paying a quadratic cost parameterized by ϕ_k . These adjustment costs are subject to an aggregate, exogenous autoregressive process κ_t , which captures changes in the efficiency of investment adjustment:

$$\ln \kappa_t = (1 - \rho_\kappa) \ln \kappa + \rho_\kappa \ln \kappa_{t-1} + \sigma_\kappa \varepsilon_{\kappa,t}. \quad (14)$$

This problem gives rise to a Tobin's Q equation, linking the investment of capital to its shadow price and its relative return.

2.3 Monetary Policy

Monetary policy operates in one of two possible regimes. In the first regime, the nominal interest rate is set according to a Taylor rule, as in [Smets and Wouters \(2007\)](#):

$$\frac{R_t}{R} = \left(\frac{R_{t-1}}{R} \right)^{\phi_r} \left(\frac{\Pi_t}{\Pi^*} \right)^{(1-\phi_r)\phi_\pi} \left(\frac{y_t}{y_t^F} \right)^{(1-\phi_r)\phi_y} \left(\frac{y_t/y_{t-1}}{y_t^F/y_{t-1}^F} \right)^{\phi_g} \exp(\sigma_R \varepsilon_{R,t}). \quad (15)$$

The nominal interest rate R_t responds to deviations in inflation Π_t from a target rate Π^* , deviations in output y_t from its flexible-price level y_t^F , and the growth rate of output relative to the growth rate of potential output, and is subject to stochastic shocks $\varepsilon_{R,t}$.

In the second regime, the nominal interest rate is at the ZLB:

$$\log(R_t) = 0. \quad (16)$$

Monetary policy can enter the ZLB regime in two ways: first, if the Taylor rule calls for negative nominal interest rates, and second, if the Fed has announced, or has previously announced, an extension of the ZLB beyond that implied by the constraint and the Taylor rule. I assume that the Fed can manipulate expectations of how the path of interest rates evolves, as in [Eggertsson and Woodford \(2003\)](#) and [Werning \(2012\)](#). In estimation, I use survey data from the New York Fed to discipline the expected duration of the zero interest rate regime during the 2009 to 2015 period.

2.4 Government

The government taxes labor income at the rate τ_t^w to fund a pay-as-you-go social security system. The period-by-period transfer paid to individuals above an eligibility age T^* depends on the accumulated pre-tax labor income of the worker, and a parameter ω governing the replacement rate of past earnings. Denote by W_t^s accumulated gross lifetime earnings, defined recursively as:

$$W_t^s = \begin{cases} w_t z^s \ell_t^s + W_{t-1}^{s-1}, & \text{if } s < T^* \\ W_{t-1}^{s-1}, & \text{if } s \geq T^*. \end{cases} \quad (17)$$

The amount ξ_t^s redistributed to an agent of age $s \geq T^*$ depends on W_t^s :

$$\xi_t^s = \omega \frac{W_t^s}{(T^* - 1)}, \quad (18)$$

where the denominator reflects the amount of time that W_t^s is accumulated over. For those younger than the eligibility age T^* , the transfer $\xi_t^s = 0$.² The budget constraint of the social security system is:

$$\sum_s n_t^s \xi_t^s = \sum_s n_t^s z^s w_t \ell_t^s \tau_t^w. \quad (19)$$

The tax rate τ_t^w adjusts to equalize social security outlays and tax revenues.

Finally, the government levies a lump-sum tax on households to pay for government expenditures g_t , which are assumed to be autoregressive and subject to stochastic shocks:

$$\ln g_t = (1 - \rho_g) \ln g + \rho_g \ln g_{t-1} + \sigma_g \varepsilon_{g,t}, \quad (20)$$

where the budget constraint of exogenous government expenditures is $g_t = n_t T_t^g$.

3 Model Approximation

To address the computational challenges arising from this model's rich sources of heterogeneity, persistence and aggregate shocks, I argue in this section that the model can be approximated by a representative agent framework with time-varying parameters that are functions of exogenous demo-

²I abstract from questions about the sustainability of pension systems in an aging society and do not allow pension funds to accumulate assets or liabilities (see [Attanasio, Kitao and Violante, 2007](#)).

graphic variables. Furthermore, because demographic changes are assumed to be perfectly foreseen, the path of these time-varying parameters are also assumed to be fully anticipated. I show how this anticipated path of time-varying parameters gives rise to a VAR representation that makes it feasible to use likelihood methods for estimation.

3.1 Derivation

The outline of the approximation follows two steps. First, the restriction that individuals can trade assets only when alive is relaxed. In the second step, we will show that when this timing assumption is relaxed, the model has a tractable aggregate representation.

Timing Assumption In the model's overlapping generations setup, individuals are born with no wealth and make period-by-period asset trades. Assume, instead, that each generation is alive at $t = 0$ and can trade claims to future consumption, and write the preferences of an individual j of age s as:

$$\sum_{t=0}^{\infty} \beta^t \left[\prod_{r=0}^s (1 - \gamma_r^r) \right] \phi_t^{j,s+t} \sum_{\sigma^t} \Pr [\sigma^t | \sigma^{t-1}] u \left[c_t^{j,s+t}(\sigma^t), \ell_t^{j,s+t}(\sigma^t) \right], \quad (21)$$

where the term $\phi_t^{j,s} = 1$ if $0 \leq s \leq T-1$, and $\phi_t^{j,s} = 0$ otherwise, indicating that individuals value utility only in the periods when they are between the ages of 0 and $T-1$. The term $\Pr [\sigma^t | \sigma^{t-1}]$ denotes the transition probability from state σ^{t-1} to σ^t . By assumption, unintentional bequests are redistributed to members of the same generation, so that individuals are insured against the idiosyncratic uncertainty associated with mortality risk. The Euler equation arising from the choice of savings in the problem where individual's maximize (21) subject to the lifetime budget constraint is:

$$\lambda_t^{j,s}(\sigma^t) = \beta \sum_{\sigma^{t+1}} \Pr [\sigma^{t+1} | \sigma^t] \lambda_{t+1}^{j,s+1}(\sigma^{t+1}) \frac{R_t(\sigma^t)}{\Pi_{t+1}(\sigma^{t+1})}, \quad (22)$$

where $\lambda_t^{j,s}(\sigma^t)$ is the Lagrange multiplier on the budget constraint. If, between any two individuals i and j , the Lagrange multipliers have the same linear relationship in the state variables, then across two periods t and t' , the ratio of the multipliers between individuals is constant:

$$\frac{\lambda_t^{j,s}(\sigma^t)}{\lambda_t^{i,s'}(\sigma^t)} = \frac{\lambda_{t'}^{j,s+t'}(\sigma^{t'})}{\lambda_{t'}^{i,s'+t'}(\sigma^{t'})} = \frac{\lambda^{i,s'}}{\lambda^{j,s}}, \quad (23)$$

where $\lambda^{j,s} = \frac{\lambda_t(\sigma^t)}{\lambda_t^{j,s}(\sigma^t)}$. The condition in Equation (23) is the same as that which arises when there are complete asset markets.

Aggregate Representation Under Equation (23), the economy's equilibrium can be found by solving the problem of a social planner that maximizes a weighted sum of individuals' utility functions. In the social planner's problem, the planner first determines how to allocate, period-by-period, aggregate consumption and aggregate labor supply between individuals. Given the optimal allocation, the planner then solves its intertemporal problem and maximizes aggregate consumption, capital, and labor supply subject to the economy's resource constraint. This approach reflects the aggregation arguments made in Constantinides (1982) and Maliar and Maliar (2003).

Assume that an individual of age s has a separable period utility function over consumption c_t^s and hours ℓ_t^s of the type $\frac{(c_t^s)^{1-\sigma}}{1-\sigma} - v^s \frac{(\ell_t^s)^{1+\varphi}}{1+\varphi}$. Under these preferences, as shown in the Appendix, there is a representative agent with preferences over aggregate consumption c_t and aggregate units of labor ℓ_t that take the form:

$$U(c_t, \ell_t) = \phi_t \frac{c_t^{1-\sigma}}{1-\sigma} - v_t \frac{\ell_t^{1+\varphi}}{1+\varphi}. \quad (24)$$

The representative agent's problem is to maximize (24) over time by choosing c_t , ℓ_t , and aggregate capital k_t subject to the economy's resource constraint and its production function $y_t = \theta_t^{1-\alpha} k_t^\alpha \ell_t^{1-\alpha}$. The relationship between the units of labor and aggregate hours is $\ell_t = A_t h_t$.

In the Appendix, I show that θ_t and A_t are the following time-varying parameters:

$$\theta_t = \sum_s n_t^s z^s, \quad \text{and} \quad A_t = \frac{\sum_s n_t^s (\hat{z}^s)^{1+1/\varphi} (v^s \lambda^s)^{-1/\varphi}}{\sum_s n_t^s (\hat{z}^s)^{1/\varphi} (v^s \lambda^s)^{-1/\varphi}}, \quad (25)$$

where the value $\hat{z}^s = z^s / \theta_t$ denotes individual s 's relative productivity and the λ^s parameters are the Pareto weights attached to an individual of age s . The shock θ_t encodes changes in output caused by variation in the size of the workforce and its composition over idiosyncratic skill levels. In particular, larger populations imply a higher θ_t , as do populations with more productive workers (or more z^s workers). The time-varying parameter A_t affects the hours needed to obtain an effective unit of labor, and is a population-weighted average of relative productivity and the disutility of providing labor. When labor supply is inelastic, θ_t and A_t together affect the labor input by $\frac{\sum_s n_t^s z^s}{\sum_s n_t^s}$: the labor input reflects the population's composition only.

The term ϕ_t inversely affects the marginal utility of consumption, and has a simple expression

mapping to the size of the population at each point in time:

$$\phi_t = \left[\sum_s n_t^s (\lambda^s)^{\frac{1}{\sigma}} \right]^\sigma. \quad (26)$$

This term scales the marginal value of wealth in the economy.

The term v_t is a time-varying parameter affecting the marginal disutility of labor:

$$v_t = \left[\sum_s n_t^s (\hat{z}^s)^{\frac{1}{\varphi}+1} (v^s \lambda^s)^{-\frac{1}{\varphi}} \right]^{-\varphi}, \quad (27)$$

so that v_t is a population-weighted average of age-specific disutilities of providing labor. The greater the relative size of the population with high disutilities of providing labor, the higher is v_t . Equating the marginal utility of consumption and the marginal disutility of labor, and substituting in for hours worked gives the labor wedge as a function of demographic changes $w_t/(\ell_t^\varphi/c_t^{-\sigma}) = v_t/\phi_t$.

Two additional trends are needed in the computations to ensure the aggregate representation will approximate the aggregate dynamics of the full lifecycle solution. The first is a gradual trend in the aggregate resource constraint to account for variation in the average mortality rate over time. The second trend is to proportional taxes used to finance the social security system. I take the path of labor income taxes from the non-stochastic perfect foresight path of the economy.

In the aggregate approximation of the model, demographics therefore affects the aggregate economy through time-varying parameters which are functions of observable population dynamics (n_t^s), the age-specific parameters of the model (z^s and v^s), and the Pareto weights that the planner attaches to each generation (λ^s). Assuming the planner equally weights each generation, the trends are straightforward to compute and do not depend on endogenous variables. In the Appendix, I verify that the aggregate approximation recovers closely the paths of the aggregate variables due to demographic shocks by comparing them to the paths of the decentralized lifecycle model under perfect foresight and demographic shocks. These results are consistent with [Ríos-Rull \(1996\)](#), who finds that the business cycle properties of large-scale, stochastic, overlapping-generations economies are similar to the properties of representative agent real business cycle models. The contribution in this paper is to additionally describe an approximation that makes likelihood estimation computationally feasible.

3.2 Solution Method

The previous section showed that under the model's approximation, demographic shocks give rise to an economy with time-varying parameters. Furthermore, since demographic changes are anticipated, this path of time-varying parameters is also anticipated. In this section, I describe the methodology used to solve for the path of the aggregated model under the anticipated path of time-varying parameters. Moreover, I describe how the methodology handles the ZLB through a regime-switching procedure.

Time-Varying Demographic Trends Let x_t be the vector of model variables (state and jump), and ε_t a vector that collects the exogenous unanticipated shocks. The linearized rational-expectations approximation of the model under time-varying parameters is:

$$\mathbf{A}_t x_t = \mathbf{C}_t + \mathbf{B}_t x_{t-1} + \mathbf{D}_t \mathbb{E}_t x_{t+1} + \mathbf{F}_t \varepsilon_t, \quad (28)$$

where \mathbf{A}_t , \mathbf{B}_t , \mathbf{C}_t , \mathbf{D}_t , and \mathbf{F}_t are time-varying matrices that encode the structural equations of the model linearized at each point in time around the steady-state corresponding to the time t structural parameters.³ A solution to the problem with anticipated time-varying parameters exists if agents expect the structural matrices to be fixed at some point in the future at values which are consistent with a time-invariant equilibrium (Kulish and Pagan, 2016). In this case, the solution has a time-varying VAR representation:

$$x_t = \mathbf{J}_t + \mathbf{Q}_t x_{t-1} + \mathbf{G}_t \varepsilon_t, \quad (29)$$

where \mathbf{J}_t , \mathbf{Q}_t , and \mathbf{G}_t are conformable matrices which are functions of the evolution of beliefs about the time-varying structural matrices \mathbf{A}_t , \mathbf{B}_t , \mathbf{C}_t , \mathbf{D}_t , and \mathbf{F}_t

$$\begin{aligned} \mathbf{Q}_t &= [\mathbf{A}_t - \mathbf{D}_t \mathbf{Q}_{t+1}]^{-1} \mathbf{B}_t \\ \mathbf{J}_t &= [\mathbf{A}_t - \mathbf{D}_t \mathbf{Q}_{t+1}]^{-1} (\mathbf{C}_t + \mathbf{D}_t \mathbf{J}_{t+1}) \\ \mathbf{G}_t &= [\mathbf{A}_t - \mathbf{D}_t \mathbf{Q}_{t+1}]^{-1} \mathbf{F}_t. \end{aligned} \quad (30)$$

This iteration is obtained by noting that, from (29), $\mathbb{E}_t x_{t+1} = \mathbf{J}_{t+1} + \mathbf{Q}_{t+1} x_t$, which is substituted into (28) and rearranged for x_t . The law of motion for the model's state variables at a time period

³One can instead linearize the model around its original steady-state, the steady-state associated with the time-varying system's final structure, or the steady-state implied by the structure at each point in time. Given the somewhat large movements in the steady-state induced by demographic changes, I use the latter approach, linearizing each set of structural matrices around the steady-state implied by that structure.

t therefore depends on the full anticipated path of the structural matrices. The final structure of the economy needed for the iteration (30) is the one that arises at the expected completion of the demographic transition and under Taylor-rule policy. Under my calibration, this final demographic structure applies from the year 2070 onwards.⁴

Zero Lower Bound To implement the occasionally-binding ZLB in the solution (29), we follow the approach of Guerrieri and Iacoviello (2015) and Jones (2017) and define two additional regimes in (28), one for when the ZLB does not bind, and one for when the ZLB binds, given the demographic parameters. If the ZLB binds, we assume that agents believe no shocks will occur in the future and iterate backwards through our model’s equilibrium conditions from the date that the ZLB is conjectured to stop binding. We then iterate on the periods that the interest rate is conjectured to be in effect until it converges, after which the solution is that of (29).⁵

4 Estimation

This section discusses how the parameters of the model are set, including the calibration of the life-cycle parameters of the model, the demographic shocks which drive the trends in the model, and the estimation of the shocks that govern the business cycle.

4.1 Assigned and Calibrated Parameters

Before estimation, a subset of the parameters and the demographic shocks are calibrated.

4.1.1 Lifecycle Parameters

The model is quarterly. Individuals begin life at 16 years of age and live for at most 80 more years, up to age 95. Full retirement is only imposed in the last period of life.⁶

I calibrate the disutility of providing labor v^s with a scaled cumulative density function of a normal distribution, so that v^s increasing in s . This specification is motivated by studies which link the

⁴This approach resembles that of Fernández-Villaverde et al. (2007) but instead of innovations driving parameter drift, the time-varying parameters are known and perfectly foreseen functions of exogenous demographic variables.

⁵See also Canova et al. (2015), Kulish and Pagan (2016), and Fernández-Villaverde et al. (2007). More generally, Jones (2017) discusses how the ZLB is a change in the structural parameters of the monetary policy rule that applies for a state-contingent period, or to stimulate the economy with forward guidance.

⁶Given the low choice of labor supply at older ages, this choice is not too important. Kulish et al. (2010) show how unanticipated changes in life expectancy can change labor supply for retired workers, as older workers with few assets return to the workforce to fund consumption during their unanticipated increase in the lifespan.

disutility of work to deteriorating health. The parameters of the function for v^s are chosen so that the labor force participation rates by age broadly match those observed in 2000. For the social security system, I set the replacement ratio of accumulated earnings λ to 46.7%, the same value that is used in [Attanasio et al. \(2007\)](#). Retirement benefits are received from age 65 on ($T^* = 49$).

I calibrate the age-productivity parameters z^s to the age-experience earnings profile. I follow [Elsby and Shapiro \(2012\)](#) in constructing the log experience-earnings profile using deflated data on full-time, full-year workers. The data is decennial Census data from 1960 to 2000, and annual American Community Survey data from 2001 to 2007.⁷ To minimize cohort effects, I pool, across years, high school dropouts, high school graduates, those with some college education, and those who have completed college or higher education.⁸ Panel A of Figure 2 plots the earnings-profile over age. The estimates imply a peak increase in wages of about 134% at age 45, before gradually declining around the age of 50: in line with the estimate of [Guvenen et al. \(2015\)](#) who find an increase in the earnings of the mean worker of 127%. There is less reliable data on the earnings of older workers, so after age 65, I calibrate the productivity of workers to decay by 20% a year.

I assess the calibration of the lifecycle parameters in Figure 2 by plotting, in Panel B, the labor force participation rate by age, in 2000, in the model and the data, and in Panel C, the age-profile of assets normalized to asset holdings at age 60, in 2013, in the model and in the Survey of Consumer Finances. The calibration implies labor force participation rates that rise when young, flatten out during an individual's prime working life, and decline rapidly around retirement ages. The lifecycle asset-profile is hump-shaped and peaks around 60 years old. Individuals borrow when young in the model and begin accumulating assets around the age of 35.

4.1.2 Mortality Profiles

Next, I calibrate the mortality probabilities of each generation during the 80 years they could possibly live, γ_t^s , to the actuarial probabilities reported by the Social Security Administration.⁹ By calibrating to these probabilities, I also match changes in the life expectancy of each generation over time, conditional on an individual reaching 16 years of age. The values used are the cohort-specific survival rates

⁷Computed off IPUMS-USA extracts. A full description is given in the Appendix.

⁸In robustness exercises reported in the Appendix, I distinguish between education groups and analyze how anticipated changes in the earnings profile map to labor supply decisions. The results indicate that the patterns of the aggregate variables are similar, suggesting that compositional changes in the population primarily drive aggregate dynamics.

⁹These probabilities were sourced from Table 7 from the Cohort Life Tables for the Social Security Area by Calendar Year, in Actuarial Study No. 120 by Felicitie C. Bell and Michael L. Miller, available at: <https://www.ssa.gov/oact/STATS/table4c6.html>. A full description is given in the Appendix.

computed for the cohort year of birth. These profiles include both observed survival rates of cohorts up to their current age, and extrapolated survival rates based on the Social Security Administration’s forecasts of life expectancy. I assume that all changes to these actuarial probabilities are exogenous and perfectly foreseen. For the initial γ_t^s profile, I use the survival probabilities reported for those born in 1900 onwards. For those cohorts born before 1900 but who are alive in 1940, I use extrapolated values of the survival probabilities.¹⁰ Under the calibration, between 1950 and 2020, life expectancy for 16 year olds increases from about 77 years to about 85 years.

4.1.3 Incoming Cohort Size Shocks

I choose anticipated shocks to the size of the incoming cohort so that the change in the observed cohort share is the same as the change in the model cohort share.¹¹ This ensures that the model captures the baby-boomer generation and imperfectly captures changes in the population distribution due to, for example, immigration. I assume that changes to the incoming population beyond 2015 decay to zero, so that the population distribution converges to the steady-state implied by the mortality profile that is constant from 2070.

I plot in Panel D of Figure 2 the median age of the population above 16 years of age implied by the calibrated mortality profiles and the incoming cohort size shocks. The profile tracks well the corresponding median age of those above 16 years of age in the data, declining from around the 1960s to around the 1980s to about 37 years of age, before steadily increasing as the baby-boomer population ages and longevity continues to rise.

4.1.4 Preference and Nominal Parameters

I calibrate the remaining parameters to values which imply steady-state capital-output ratios that align with those in the Bureau of Labor Studies’ Multifactor Productivity (BLS-MFP) program (see Fernald, 2015).¹² I set capital to depreciate by $\delta = 10.6\%$ a year. The capital share α is set to $1/3$, the average of the capital share reported by Fernald (2015) over 1948 to 2015. The intertemporal elasticity of substitution σ is set to 1, and the inverse Frisch elasticity of labor supply φ is set to 2, in line with the estimates of Reichling and Whalen (2012) and with the analysis of Rios-Rull et al. (2012).

¹⁰Because the survival probabilities are low for those years, the results are robust to alternative specifications and are not important for the model outcomes beyond 1970.

¹¹Choosing initial population shocks to matching the *changes* is necessary because the model is initialized at the 1940 steady-state and matching the actual cohort sizes would imply very large and counterfactual initial population shocks.

¹²These observed capital-output ratios vary between 2 and 2.7 over the period 1950 to 2013. A description of the BLS-MFP dataset used is given in the Appendix.

Both values are also in the range considered by [Auerbach and Kotlikoff \(1987\)](#) in a computational overlapping generations model. The quarterly discount factor β is set to 0.99875. Together, these parameters imply a capital-output ratio in 2000 of about 2.7, which is the observed capital-output ratio in the BLS-MFP in 2000.

I calibrate a small set of the parameters describing the nominal side of the economy to values commonly used in the literature. The steady-state value of ξ is set at 8, which implies a steady-state markup over marginal costs of $\xi/(\xi - 1) = 14\%$. The parameter governing the quadratic cost of capital adjustment ϕ_k is set to 40. The slope of the Phillips curve, ϵ_p , and therefore the quadratic cost of price adjustment ϕ_p is estimated in the next section. I estimate this parameter because it is key for determining the behavior of inflation and therefore the real interest rate, particularly in response to changes in the expected ZLB duration which are due to forward guidance announcements. The annual inflation target Π^* is set to 2.2%. Finally, I use the values that [Smets and Wouters \(2007\)](#) estimate for the Taylor rule parameters, which at their posterior mode are: $\phi_r = 0.81$, $\phi_\pi = 2.03$, $\phi_y = 0.08$, and $\phi_g = 0.22$.

4.2 Bayesian Estimation

I next use Bayesian likelihood techniques to estimate the slope of the Phillips curve, trend growth, and the model's shocks that drive business cycle fluctuations around the demographic trend.

4.2.1 Quarterly Data

The solution expressed in equation (29) has a state-space representation, allowing me to adapt Bayesian likelihood methods to estimate the remaining parameters of the model. The quarterly data used are:

$$\text{Data} = \left\{ \log \left(\frac{y_t}{y_{t-1}} \right), \log \left(\frac{c_t}{c_{t-1}} \right), \log \left(\frac{i_t}{i_{t-1}} \right), \log \Pi_t, \log \tilde{R}_t, T_t \right\}_t, \quad (31)$$

over the time period 1984Q1 to 2015Q1. I use, as observables, the growth rate of output per capita, of consumption per capita, of investment per capita, the GDP deflator, and the Fed Funds rate, and follow [Smets and Wouters \(2007\)](#) in constructing these series.¹³ The nominal interest rate is no longer an observable when the ZLB binds between 2009Q1 and 2015Q1. I implement this with a time-varying observation equation in the state-space representation of the model (see [Kulish et al., 2017](#)). The

¹³The Appendix provides more details of the data series used in estimation.

sequence of expected durations of the ZLB, \tilde{T}_t , between 2009 and 2015 are taken from the Blue Chip Financial Forecasts survey from 2009 to 2010 and the New York Federal Reserve’s Survey of Primary Dealers from 2011 to 2015.

4.2.2 Parameter Estimates

Table 1 reports moments of the prior and posterior distributions. The priors are diffuse. The prior distribution over the slope of the Phillips curve is wide and allows for a high degree of price flexibility.¹⁴ The prior distribution for the trend rate of growth z is also wide, with the 10th (90th) percentiles implying annual trend growth of about 1% (3%). I use a Markov Chain Monte Carlo approach to characterize the parameters’ posterior distributions, computing two independent chains of 150,000 draws. The Appendix provides additional details of the estimation and an analysis of the convergence of the chains.

First, I find that prices are quite inflexible, with the posterior estimate of the slope of the Phillips curve ϵ_p centered tightly around 0.01. At the posterior mode, this value translates into a Calvo probability of price adjustment every quarter of about 10%, consistent with the estimates in [Smets and Wouters \(2007\)](#), but a little more flexible than the estimates of [Del Negro et al. \(2015\)](#) and [Aruoba et al. \(2017\)](#).

Next, I find that the modal estimate of trend growth implies an annual growth rate of around 1.35%. This is about 0.4% annual percentage points less than the estimate implied by [Smets and Wouters \(2007\)](#), in part because my estimation sample includes a period of lower average growth, but also because demographic trends alone account for some of the growth over this period, as discussed in the next section.

The remaining parameters are those governing the persistence and size of the shock processes. To interpret these, I report, in Table 2, the unconditional and one-year forecast error variance decomposition of the observable variables (and wages). These decompositions reveal how important each shock is in driving the observable variables around the demographic trend. At the one-year horizon, monetary policy shocks, exogenous government spending shocks, and investment shocks together make up about half of the forecast error variance for output, while at a longer horizon, about 70 percent of the forecast error variance is caused by markup shocks. About one-third of the forecast error variance of the

¹⁴Translating the quadratic price adjustment cost into a Calvo price-reset probability, the 10th percentile of the prior distribution for the slope implies a quarterly Calvo reset probability of 13%, while the 90th percentile of the prior distribution of the slope implies a Calvo reset probability of 34%.

Fed Funds rate is driven by preference shocks. As discussed in the next section under counterfactual simulations, these shocks are largely accommodated by monetary policy, but can be very contractionary when monetary policy is constrained.

Monetary policy shocks do not affect output, consumption, and investment much, comprising around 5% of the forecast error variance at the one-year horizon, and less than 1% at the infinite horizon, which is close to the fraction reported by [Smets and Wouters \(2007\)](#). Monetary policy shocks instead account for about 23% of the forecast error variance of the Fed Funds rate at short horizons. Finally, consistent with the results in [Justiniano et al. \(2011\)](#), investment shocks are very important for output and investment at both the one-year and infinite horizons, explaining 27% and 50% of the forecast error variance of output and investment at the one-year horizon, and 18% and 30% of the forecast error variance of output and investment at the infinite horizon.

I gauge the importance of including demographic shocks in the estimation by reporting, in the Appendix, the posterior distributions of the estimated parameters and the forecast error variance decomposition of the observables in an estimation where demographics are held constant from 1984 on. I find that the estimate of annual trend growth is higher by 0.4%, since demographic trends contribute positively to growth, as discussed in more detail in the next section. I also find a much more important role for technology shocks in explaining output and consumption when demographics are not explicitly taken into account. This suggests that demographic trends manifest themselves in ways that reflect highly persistent changes in productivity.

5 Demographics and the Business Cycle

I next study, using the estimated model, the role that demographic changes have had in explaining the decline in log output relative to its pre-crisis trend. I first show that demographic shocks alone are responsible for about one-third of the decline in output. On top of this, demographic changes have an additional nonlinear effect by causing the ZLB to bind between 2009 and 2015. I show this by computing a counterfactual holding demographics constant from 1984. After removing the forward guidance response of the Fed to a binding ZLB, I find that output would have fallen by an additional 2% relative to trend, primarily between 2011 and mid-2013.

5.1 Effect of Demographics

Effect of Demographic Shocks Alone First, I discuss how demographic shocks alone affect the economy’s key variables, plotted in Figure 3. Starting in 1984, I turn off all shocks except those to fertility and mortality rates. Panel A shows that, between 1990 and 2015, the Fed Funds rate declines by about 2 percentage points. Panel B plots an index of log output and illustrates how demographic changes cause a slowdown in output growth relative to log output’s 1984 to 2007 trend. In Panel C of Figure 3, I show that the real interest rate is expected to fall by about 1 percentage point – driven by changes in the capital-output ratio – while in Panel D, demographic shocks cause a decline in employment of about 2.5% between 1990 and 2015.

Regarding the implications of demographics alone for growth, from 1980 to 2015, my model predicts growth falls by about 1.25 percentage points. There are three main channels through which output growth can change over time because of changing demographics. Workers can supply more hours, affecting both output and aggregate labor. There are also changes in physical capital, as individuals save and consume out of accumulated savings in retirement. Third, the *quality* of labor can change; namely, changes in the distribution of workers resulting from demographic changes alters the average skill-level of the workforce, which shows up in a decomposition of productivity growth as fluctuations in the average productivity of labor (Fernald, 2015).

I decompose the model’s predictions for output growth and labor productivity growth into their component parts and show that accelerating capital accumulation increases the growth rate of both labor productivity and total output up to 1995, after which the growth rate starts to decline.¹⁵ The change in labor supply has a large negative effect on productivity growth, but a positive effect on total growth, when the baby boomer cohorts enter the labor force around 1960. A key component of both labor productivity and total growth is the change in the average skill level of the workforce caused by the interaction of a changing composition of the workforce with the age-productivity profile. The decomposition implies that the contribution of the change in average labor quality to the growth rate of output and output per worker peaks around 1990, adding roughly 0.3 percentage points to total growth and productivity growth. The contribution of labor quality becomes a drag on productivity growth in 2000 as a large fraction of workers reach the peak of the age-productivity profile, exhausting the potential for further growth in average human capital across the workforce. This force is forecast to depress productivity growth until 2030. In total, I find that demographics will be a drag on output

¹⁵This decomposition is presented in more detail in the Appendix.

growth through to 2070.

In Figure 4, I explore the trends generated by demographics only in more detail by plotting first, in Panel A, the employment population ratio in the model and the data. Demographics alone capture the dynamics of the aggregate employment-population ratio well under the calibration of the lifecycle parameters, which generated age-specific labor force participation rates that are consistent with those observed. The labor force participation rate declines in the model at a pace that is roughly as fast as that observed and is predicted to continue to fall by a further 4 percentage points from 2020 to 2040. This result is driven by the compositional changes in the workforce towards workers with lower participation rates (as shown in Panel B of Figure 2), a point supported empirically by [Aaronson et al. \(2014\)](#).¹⁶

Demographic trends have important implications for the capital-output ratio, plotted in the second panel of Figure 4. As life expectancy rises and mortality rates fall, aggregate savings increases to fund longer expected retirements. As a result, the capital-output ratio increases and the marginal product of capital and real interest rate fall. In addition, the aging of the baby boomer cohorts generates an increase, and then decrease, in the path for the real interest rate around the secular decline implied by increasing longevity. The oscillation is driven by changes in the relative size and composition of the workforce. The workforce is relatively young as the baby-boomers enter the labor market in the 1960s to 1980s, so that aggregate hours supplied is high relative to capital, thereby increasing the marginal return to capital. As the baby-boomer cohort ages and accumulate savings for retirement, the marginal return to capital and the real interest rate decline. This decline is then reinforced by the withdrawal of the baby-boomer cohort from the labor market, depressing the marginal return to capital, which stays low beyond 2030 (see also [Carvalho et al., 2016](#)).

These results on the implications of demographics for macroeconomic trends are consistent with the findings of other studies. Using a calibrated overlapping generations model, [Gagnon et al. \(2016\)](#) find that demographic trends generate a decline in the growth rate from around 2.3 percentage points in 1980 to just under 0.5 percentage points by 2030. They find, similar to what is predicted by my model, that much of the decline is due to declining fertility and the associated exit of the baby boomer

¹⁶Furthermore, to decompose the contribution of the mechanical effect of an aging population to the decline in the labor input, we can compare (i) the labor input predicted by the model in our baseline exercise where labor endogenously responds to wages, to (ii) the labor input predicted by the model when labor supply is inelastic ($\varphi \rightarrow \infty$). In this comparison, almost all of the forecasted decline in the labor input is due to the mechanical effect of demographic changes; the endogenous response of labor in my model to demographic changes mitigates the 8 percent decline (relative to 2015) by only 2 percentage points.

generations from the labor force.¹⁷ In an empirical study, [Aksoy et al. \(2019\)](#) predict an average decline in annual output growth rates across OECD countries of $1\frac{1}{4}$ percentage points between 2010 and 2030. Regarding real interest rates, [Gagnon et al. \(2016\)](#) find a peak in the real rate between 1975 and 1985 of around 1.7%, and a decline of about 1.4 to 1.5 percentage points by 2030, close to my model’s predictions. In an empirical study, [Johannsen and Mertens \(2016\)](#) find a decline in the real rate of about $\frac{3}{4}$ percentage points between 1985 and 2015, while [Eggertsson et al. \(2019\)](#) forecast a larger decline in their model implied interest rate due to demographics of around 3 percentage points between the 1980s and 2030.

Decline in Real Interest Rate and the Zero Lower Bound Next, I examine the nonlinear interactions between the decline in the real interest rate and the ZLB on the Fed Funds rate. In Figure 5, I plot the response of the Fed Funds rate and output to a large negative investment shock under two different demographic states – one initialized at the steady-state associated with the 1990 demographic profile, and the other for the steady-state associated with the 2008 demographic profile. The Fed Funds rate is higher in steady-state for the 1990 demographics, when the population is younger the supply of savings is relatively lower. The two responses illustrate how the same shock can have very different implications for the economy because of the binding ZLB.¹⁸ Initialized at the 2008 steady-state, the shock is large enough to cause the ZLB to bind for about two years, with output falling by an additional 2 percentage points on impact. These impulse responses show that the decline in the real interest rate can be quantitatively important in the presence of the ZLB, as discussed next.

Holding Demographics Constant from 1984 Here, I examine whether demographic trends were responsible for the Fed Funds rate hitting the ZLB between 2009 and 2015. To answer this, I hold the demographic profile constant at its 1984 state, which was the first year that quarterly data is used in the Bayesian estimation, and construct a counterfactual using the estimated structural shocks. As illustrated in Panel A of Figure 6, had the population not aged between 1984 and 2015, the Fed Funds rate would have remained above zero in the aftermath of the Great Recession, falling to, at its lowest point, 0.5% in annual terms in 2012Q4.¹⁹ The declining Fed Funds rate caused by demographic changes implies that, going forward, the ZLB is likely to be visited more frequently. One potential

¹⁷I present decompositions of these contributing factors in the Appendix.

¹⁸Due to the time-varying structural changes induced by demographics, there are also small differences in the propagation of shocks unrelated to the binding lower bound. I present additional results of such differences in the Appendix.

¹⁹This period corresponds to the time of the lowest 10-year Treasury yields, which were around 1.6%.

policy response would be to raise the inflation target, thereby raising the steady-state nominal interest rate. In simulations, I find that the annual inflation target would need to be raised to about 3.5% to obtain approximately the same ex-ante distribution of expected ZLB episodes in 2015 as would arise under an inflation target of 2% and when demographics are set to their 1984 state.

Panel B shows the effect of unchanged demographics for output. Interestingly, the counterfactual response shows that output growth would have been lower. The reason for this is that the composition of the workforce in 1984 is skewed towards younger workers and is therefore favorable to productivity growth between 1984 and the mid-2000s, as younger cohorts move up the age-productivity profile, generating productivity growth. Panel D shows that holding the demographic profile fixed at its 1984 value, the employment gap would have been closed just prior to the recession, and in 2009 would have been about 2 percentage points narrower, consistent with the results presented in Figure 3 that demographic shocks have caused a secular decline in the employment-population ratio.

5.2 Effect of Forward Guidance

The previous section showed how demographic trends over the past 25 years caused a decline in real and nominal interest rates, thereby making the ZLB a constraint on monetary policy. I explore how much of a constraint the ZLB was, quantitatively, by constructing a counterfactual simulation of the economy in which the Fed acts passively in response to shocks that cause the ZLB to bind. In this simulation, the expected ZLB duration adjusts in response to the shock only. By contrast, in the estimation, I fix the expected ZLB duration to those observed in survey data, which allows the Fed to extend the ZLB duration beyond the duration implied by the shocks themselves (see also [Campbell et al., 2012](#); [Jones, 2017](#)). The counterfactual simulation therefore provides a measure of the degree to which the ZLB was a binding constraint, absent explicit forward guidance policies.

Figure 7 illustrates the counterfactual path of the economy without forward guidance. Panel C shows that inflation would have been lower by about 1 to 2 percentage points between 2009 and 2015, and Panel D illustrates that employment would have been lower by a further 0.7 percentage points in 2012Q2. Cumulatively, over the 2009 to 2015 period, employment would have been lower by 7.6%.

Next, I compute the ZLB durations implied by the shocks alone, which provides a measure of how stimulatory forward guidance is.²⁰ I find some degree of forward guidance stimulus every quarter

²⁰I calculate the ZLB durations implied by the structural shocks, using a method described in [Jones \(2017\)](#). The difference between those computed endogenous durations and the durations used in the estimation is the contribution of forward guidance, or the extension of the ZLB regime that, together with the structural shocks, will generate the observed series. The decomposed ZLB durations are plotted in the Appendix.

between 2009 and 2015, but the strongest forward guidance announcements occur between 2011Q3 and 2013Q3, when the forward guidance component of the total duration is estimated to be between 8 and 9 quarters. This period corresponds to low yields on long-term Treasuries and the explicit calendar-based targets announced by the Fed. These results are also consistent with the findings in [Swanson and Williams \(2014\)](#), who show that between 2009 and 2011, long-term yields were relatively unconstrained, and that after 2011, long-term yields tightened significantly towards their lower bounds; consistent with the Fed announcing expansive unconventional monetary policies. In particular, around 2011, the Fed announced its “to mid-2013” guidance announcement, the first of many subsequent calendar-based extensions of the lower bound regime.²¹

5.3 Output Since the Great Recession

In this section, I put together the results from the previous two sections to study how the model decomposes the decline in log output relative to its long-run linear trend. First, Figure 8 plots the difference between output and its long-run trend relative to 2008. The data show the severity of the slowdown in output, with the gap between the data and its long-run trend widening between 2008 and 2015 to 12% by 2015. Demographics alone account for a third – 4 percentage points – of the gap in 2015. Without expansionary forward guidance policy, the gap would have been larger by, at most, 2 percentage points, primarily between 2011 and 2013.

Because of forward guidance policy and the nonlinearities associated with the ZLB, there is no linear decomposition of the remaining gap between output and its demographic path into the contribution of individual shocks.²² For this reason, I explore how each estimated shock affected output in two cases: first, where the ZLB duration adjusts in response to changes in shocks, and second, where the sequence of ZLB durations are held constant at their observed values.

Figure 9a plots the contribution of each shock to output, computed in the following way. First, I compute counterfactual paths of output when each shock is set to zero, respectively, and the Fed Funds rate responds endogenously. I then evaluate the contribution of each shock as the difference between the observed output series and the counterfactual output series. The results suggest that investment efficiency shocks primarily explain why output is low relative to trend, and have a persistent negative effect on output through to 2015. These shocks, as emphasized by [Justiniano et al. \(2011\)](#), capture

²¹For example, in the FOMC press release, August 9 2011, the FOMC announced: “The Committee currently anticipates that economic conditions – including low rates of resource utilization and a subdued outlook for inflation over the medium run – are likely to warrant exceptionally low levels for the federal funds rate at least through mid-2013.”

²²The shocks are presented in the Appendix.

disruptions in the intermediation between savings and investment. They therefore capture the impact that financial factors have on aggregate variables which, as shown by [Jermann and Quadrini \(2012\)](#) and [Del Negro et al. \(2015\)](#), are key to explaining the decline in output around 2008-09.²³ The results also suggest that government spending shocks kept output from falling even further between mid-2009 and 2012. Discount factor shocks have a small effect on output when monetary policy is unconstrained and able to respond. This observation is consistent with the variance decompositions of Table 2, where preference shocks explain about a third of the forecast error variance of the Fed Funds rate, but a small fraction of the variance of output, consumption, and investment.

Finally, in Figure 9b, I compute the contribution of each shock holding fixed the sequence of ZLB durations between 2009 and 2015. This exercise magnifies the role of each shock, since the Fed does not endogenously react when that shock is removed. Government spending shocks have a larger marginal impact because these shocks stimulate inflation and lower real interest rates more when nominal interest rates are fixed (see also [Christiano et al., 2011](#)). Similarly, markup shocks have a positive effect on output when the interest rate is fixed in 2012, because they raise inflation and lower real interest rates. The estimated preference shocks have a substantial contractionary effect on output when the interest rate is fixed. These shocks have a large effect on the real interest rate when they are not accommodated by monetary policy, thereby affecting consumption and output.

5.4 Robustness

I conduct a number of experiments to verify that the aggregate trend predictions from the model under the expected demographic changes are robust to alternative specifications of the model. The results of each experiment are presented in the Appendix but discussed briefly here.

Borrowing Constraints I first check that the model’s predictions hold when individuals face a constraint restricting their borrowing early in life. With borrowing constraints, there are more savings, pushing up the capital-output ratio. As a consequence, the real interest rate is lower than in the baseline model. The magnitude of the fluctuations of the real interest rate, the participation rate and output growth are very similar to the baseline model.

²³The shocks to the efficiency of investment – shocks to κ_t in (5) – correlate with the spread between the yields of Baa Corporate debt and 10 year Treasury bonds, and spike at the same time as the spike in the spread in 2008.

Time-Varying Productivity Profile The second robustness check is to adopt time-varying productivity profiles to account for a possible flattening of productivity profiles over time. Such a flattening can affect the accumulation of human capital and can impact aggregate productivity measures in two ways: first, by a growth effect, by lowering the potential for new workers to accumulate human capital, and second, by a level effect, by affecting the productivity level that individuals enter the workforce on. I calibrate the age-productivity profiles by recomputing for each cross-sectional sample, the profile and then interpolating between those points in time. The overall pattern of aggregate labor productivity is much the same as the baseline model, although the magnitude of the amplitude of the change in labor productivity growth is smaller, with demographics contributing the most to labor productivity growth in 1980 rather than in 1990 (as in the baseline results).

Female Labor Force Participation and Multiple Skill Types From 1985 on, the baseline predictions for the participation rate, aggregate labor productivity growth and the real interest rate are largely unaffected when the age-productivity and labor disutility profiles are calibrated to match female age-earnings profiles and female labor force participation rates from the 1940s to 1990s, after which female labor force participation is roughly constant. As a final point of comparison, I verify that the directions of the aggregate predictions are robust to a calibration where an additional source of heterogeneity is modeled—where there are two types of workers, low or high skilled, where low skill workers are those with less than college education. These robustness exercises emphasize how the important demographic dynamics are captured primarily through changes in the size of the population.

6 Conclusions

This paper studies why the level of US output remains significantly below its pre-crisis trend after the Great recession. I use a New Keynesian model with demographic shocks and the ZLB to show that declining mortality rates and changes to the age population composition can generate long-run trends that match the low frequency movement of output growth, productivity, the real interest rate, and the employment-population ratio.

I estimate the model using Bayesian likelihood methods and accounting for ZLB, forward guidance, and the demographic transition. With the estimated model, I find that the ZLB would not have been a binding constraint between 2009 and 2015 had there been no demographic shocks from 1984 and the population had not aged. I find that demographic shocks alone are responsible for about one-third

of the decline in output relative to its pre-crisis trend by 2015. Furthermore, my results suggest that absent any forward guidance policy used by the Fed, the ZLB would have caused output to fall by an additional 2 percentage points between 2011 and 2013.

The shocks themselves can have highly nonlinear effects on output and consumption. I assess the contribution of each of the estimated shocks to the decline in output since the Great Recession, and find an important role for investment shocks in causing output to fall. These shocks proxy for financial disturbances as they capture disruptions in the intermediation between savings and investment.

The results illustrate the importance of demographics as a major driver of macroeconomic trends over time. Further research could focus on how demographic trends interact with the housing market or with the efficacy of fiscal policy. It would also be interesting to model a more detailed financial sector to study the interactions between demographic changes and financial frictions, particularly as borrowing constraints may bind differently across cohorts.

References

- Aaronson, S., Cajner, T., Fallick, B., Galbis-Reig, F., Smith, C., Wascher, W., 2014. Labor Force Participation: Recent Developments and Future Prospects. *Brookings Papers on Economic Activity* Fall, 197–255.
- Aksoy, Y., Basso, H.S., Grasl, T., Smith, R.P., 2019. Demographic Structure and Macroeconomic Trends. *American Economic Journal: Macroeconomics* 11, 193–222.
- Aruoba, S., Cuba-Borda, P., Schorfheide, F., 2017. Macroeconomic Dynamics Near the ZLB: A Tale of Two Countries. *Review of Economic Studies* .
- Attanasio, O., Kitao, S., Violante, G.L., 2007. Global Demographic Trends and Social Security Reform. *Journal of Monetary Economics* 54, 144–198.
- Auerbach, A.J., Kotlikoff, L.J., 1987. *Dynamic Fiscal Policy*. Cambridge University Press.
- Campbell, J.R., Evans, C.L., Fisher, J.D., Justiniano, A., Calomiris, C.W., Woodford, M., 2012. Macroeconomic Effects of Federal Reserve Forward Guidance. *Brookings Papers on Economic Activity* Spring, 1–80.
- Canova, F., Ferroni, F., Matthes, C., 2015. Approximating Time Varying Structural Models with Time Invariant Structures.
- Carvalho, C., Ferrero, A., Nechio, F., 2016. Demographics and Real Interest Rates: Inspecting the Mechanism. *European Economic Review* 88, 208–226.

- Christiano, L., Eichenbaum, M., Rebelo, S., 2011. When Is the Government Spending Multiplier Large? *Journal of Political Economy* 119, 78 – 121.
- Constantinides, G.M., 1982. Intertemporal asset pricing with heterogeneous consumers and without demand aggregation. *The Journal of Business* 55, 253–267.
- Del Negro, M., Giannoni, M., Schorfheide, F., 2015. Inflation in the great recession and new keynesian models. *American Economic Journal: Macroeconomics* 7, 168–196.
- Eggertsson, G.B., Mehrotra, N.R., Robbins, J.A., 2019. A Model of Secular Stagnation: Theory and Quantitative Evaluation. *American Economic Journal: Macroeconomics* .
- Eggertsson, G.B., Woodford, M., 2003. The Zero Bound on Interest Rates and Optimal Monetary Policy. *Brookings Papers on Economic Activity* 1, 139–233.
- Elsby, M.W.L., Shapiro, M.D., 2012. Why Does Trend Growth Affect Equilibrium Employment: A New Explanation of an Old Puzzle. *American Economic Review* 102, 1378–1413.
- Fernald, J., 2015. Productivity and Potential Output before, during, and after the Great Recession. *NBER Macroeconomics Annual* 29, 1–51.
- Fernández-Villaverde, J., Gordon, G., Guerrón-Quintana, P.A., Rubio-Ramírez, J., 2007. How Structural Are Structural Parameters? *NBER Macroeconomics Annual* 22, 83–132.
- Feyrer, J., 2007. Demographics and Productivity. *The Review of Economics and Statistics* 89, 100–109.
- Gagnon, E., Johannsen, B.K., Lopez-Salido, D., 2016. Understanding the New Normal: The Role of Demographics. Working Paper. Board of Governors of the Federal Reserve System.
- Guerrieri, L., Iacoviello, M., 2015. Occbin: A Toolkit to Solve Models with Occasionally Binding Constraints Easily. *Journal of Monetary Economics* 70, 22–38.
- Guvonen, F., Karahan, F., Ozkan, S., Song, J., 2015. What Do Data on Millions of U.S. Workers Reveal about Life-Cycle Earnings Risk? Working Paper 20913. National Bureau of Economic Research.
- Jaimovich, N., Siu, H., 2012. The Trend is the Cycle: Job Polarization and Jobless Recoveries.
- Jermann, U., Quadrini, V., 2012. Macroeconomic effects of financial shocks. *American Economic Review* 102, 238–71. URL: <http://www.aeaweb.org/articles.php?doi=10.1257/aer.102.1.238>, doi:10.1257/aer.102.1.238.
- Johannsen, B.K., Mertens, E., 2016. A Time Series Model of Interest Rates With the Effective Lower Bound. Technical Report 2016-033. Washington: Board of Governors of the Federal Reserve System.

- Jones, C., 2017. Unanticipated Shocks and Forward Guidance at the Zero Lower Bound. NYU.
- Justiniano, A., Primiceri, G.E., Tambalotti, A., 2011. Investment Shocks and the Relative Price of Investment. *Review of Economic Dynamics* 14, 102–121.
- Kulish, M., Kent, C., Smith, K., 2010. Aging, Retirement, and Savings: A General Equilibrium Analysis. *The BE Journal of Macroeconomics* 10.
- Kulish, M., Morley, J., Robinson, T., 2017. Estimating DSGE Models with Zero Interest Rate Policy. *Journal of Monetary Economics* 88, 35–49.
- Kulish, M., Pagan, A., 2016. Estimation and Solution of Models with Expectations and Structural Changes. *Journal of Applied Econometrics* .
- Maliar, L., Maliar, S., 2003. The Representative Consumer in the Neoclassical Growth Model with Idiosyncratic Shocks. *Review of Economic Dynamics* 6, 362–380.
- Reichling, F., Whalen, C., 2012. Review of Estimates of the Frisch Elasticity of Labor Supply. Working Paper 13. Congressional Budget Office.
- Ríos-Rull, J.V., 1996. Life-cycle economies and aggregate fluctuations. *Review of Economic Studies* 63, 465–489.
- Rios-Rull, J.V., Schorfheide, F., Fuentes-Albero, C., Kryshko, M., Santaaulalia-Llopis, R., 2012. Methods Versus Substance: Measuring the Effects of Technology Shocks 59, 826–846.
- Smets, F., Wouters, R., 2007. Shocks and Frictions in US Business Cycles: A Bayesian DSGE Approach. *American Economic Review* 96, 586–606.
- Summers, L.H., 2014. U.S. Economic Prospects: Secular Stagnation, Hysteresis, and the Zero Lower Bound. *Business Economics* 49, 65–73.
- Swanson, E., Williams, J., 2014. Measuring the Effect of the Zero Lower Bound on Medium- and Long-term Interest Rates. *American Economic Review* 104, 3154–3185.
- Werning, I., 2012. Managing a Liquidity Trap: Monetary and Fiscal Policy.
- Wong, A., 2015. Population Aging and the Transmission of Monetary Policy to Consumption.

Table 1: Estimated Parameters

Parameter	Prior				Posterior			
	Dist	Median	10%	90%	Mode	Median	10%	90%
ϵ_p	U	0.1	0.0	0.2	0.01	0.01	0.01	0.02
$400 \times (z - 1)$	N	2.0	1.1	3.0	1.34	1.35	1.24	1.46
ρ_χ	B	0.5	0.3	0.7	0.94	0.94	0.93	0.95
ρ_μ	B	0.5	0.3	0.7	0.76	0.73	0.55	0.83
ρ_θ	B	0.5	0.3	0.7	0.96	0.96	0.95	0.97
ρ_g	B	0.5	0.3	0.7	0.95	0.95	0.94	0.97
ρ_κ	B	0.5	0.3	0.7	0.94	0.94	0.93	0.95
$100 \times \sigma_\chi$	IG	1.2	0.5	3.7	2.15	2.18	1.95	2.47
$100 \times \sigma_\mu$	IG	1.2	0.5	3.7	0.48	0.52	0.33	0.73
$100 \times \sigma_\theta$	IG	1.2	0.5	3.7	3.51	3.56	3.20	4.03
$100 \times \sigma_i$	IG	1.2	0.5	3.7	0.16	0.16	0.15	0.18
$100 \times \sigma_g$	IG	1.2	0.5	3.7	1.05	1.06	0.98	1.15
$100 \times \sigma_\kappa$	IG	1.2	0.5	3.7	1.00	1.04	0.86	1.25

Table 2: Variance Decomposition Due to Shocks, %

Shock Variable	Preference	Technology	Markup	Policy	Government	Investment
A. Conditional, 4 Quarter Ahead						
Fed Funds Rate	28.5	7.8	5.8	22.9	8.3	26.8
Inflation	21.9	2.4	42.1	8.9	6.4	18.4
Wages	6.8	2.6	74.8	10.5	0.6	4.8
Output	0.8	0.3	47.7	4.3	19.6	27.4
Consumption	28.0	0.7	26.7	6.8	26.7	11.2
Investment	7.2	0.1	38.8	2.3	0.8	50.8
B. Unconditional						
Fed Funds Rate	34.6	1.8	28.7	5.6	9.7	19.6
Inflation	20.5	1.1	55.7	4.2	5.8	12.7
Wages	1.7	0.4	93.0	2.0	0.2	2.7
Output	4.2	0.0	70.9	0.7	6.2	18.0
Consumption	7.0	0.1	70.0	1.0	13.6	8.4
Investment	8.6	0.0	59.1	0.5	1.2	30.5

Figure 1: Macroeconomic Trends


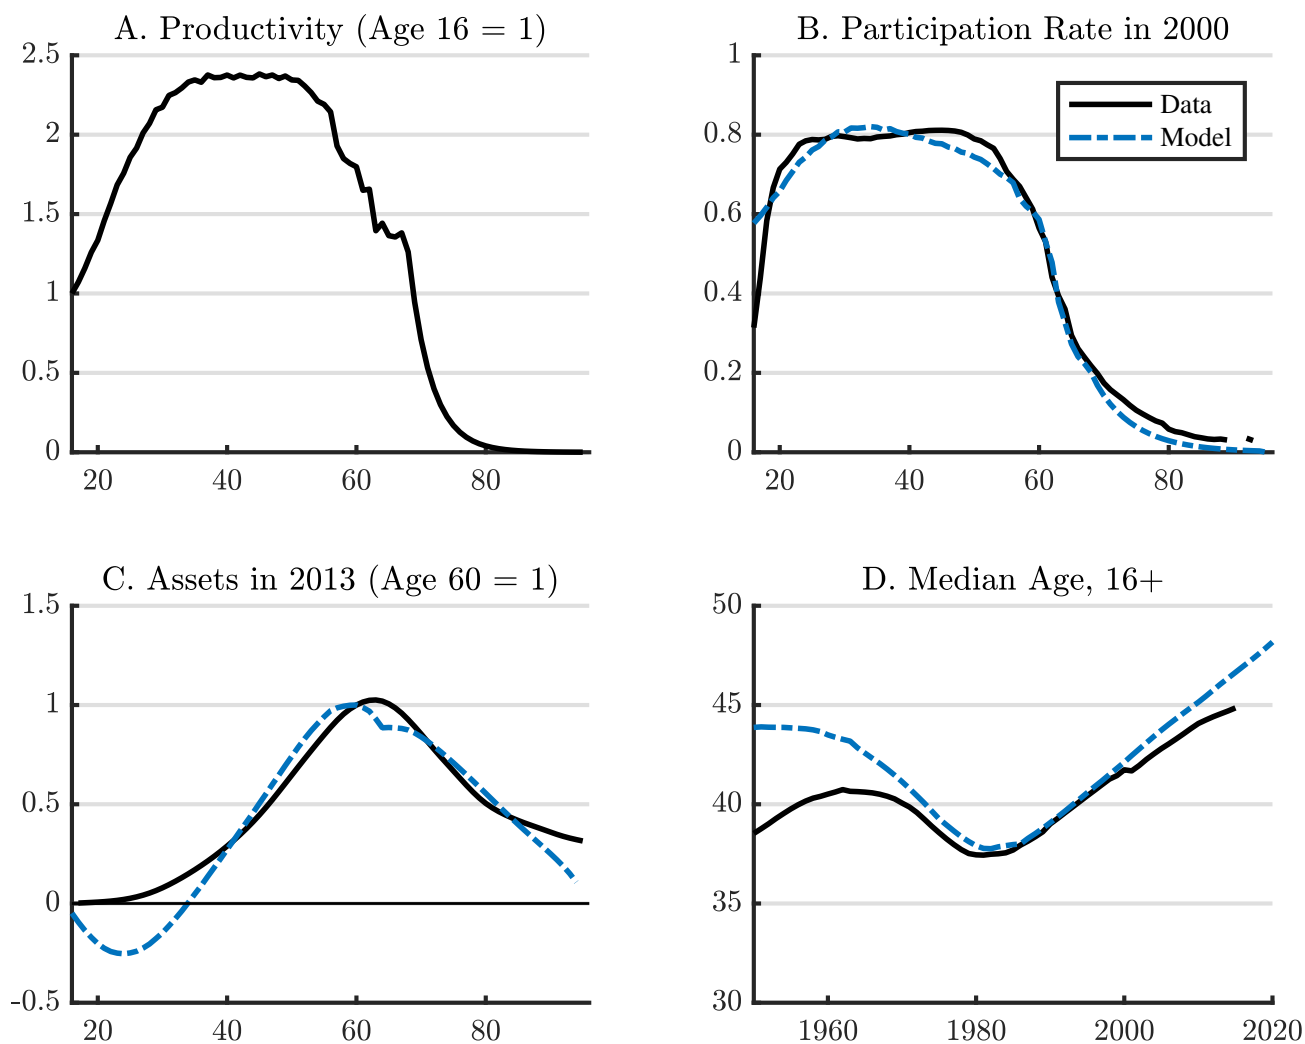


Figure 2: Individual Profiles



Notes: The data for the asset profile is taken from the Survey of Consumer Finances for 2013, HP-filtered, and normalized to the values of assets when 60. The productivity profile is computed from pooled Census and American Community Survey datasets. All details are in the Appendix.

Figure 3: Path of Variables, Demographics Only

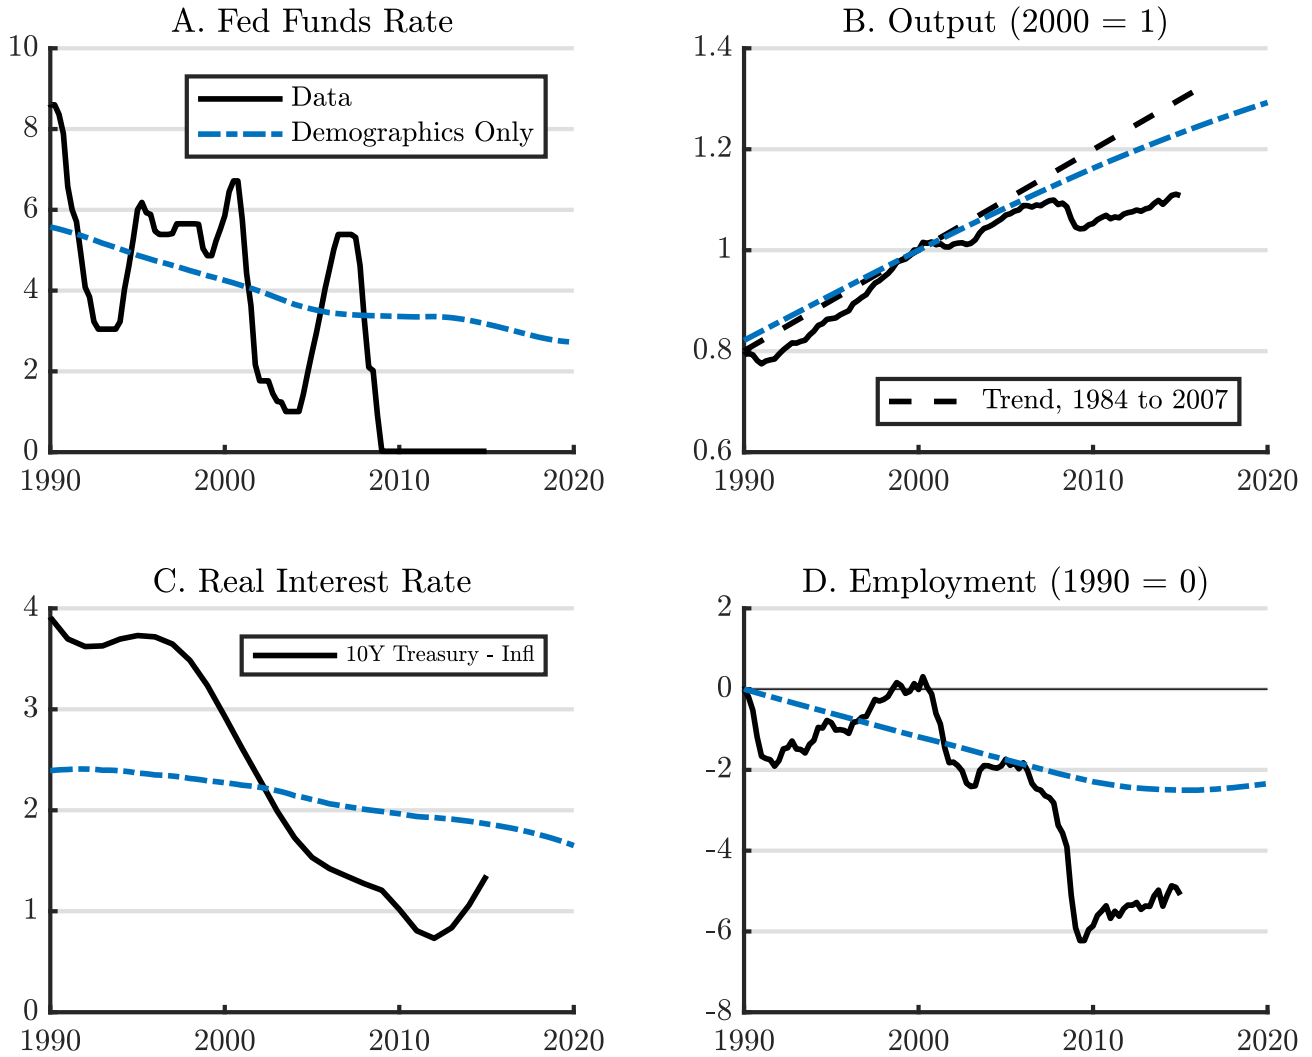



Figure 4: Model Trends



Notes: The data for the capital-output ratio is extracted from the BLS's Multifactor Productivity program. All details are in the Appendix.

Figure 5: Impulse Response to Large Investment Shock




Figure 6: Path of Variables, Demographics Fixed from 1984

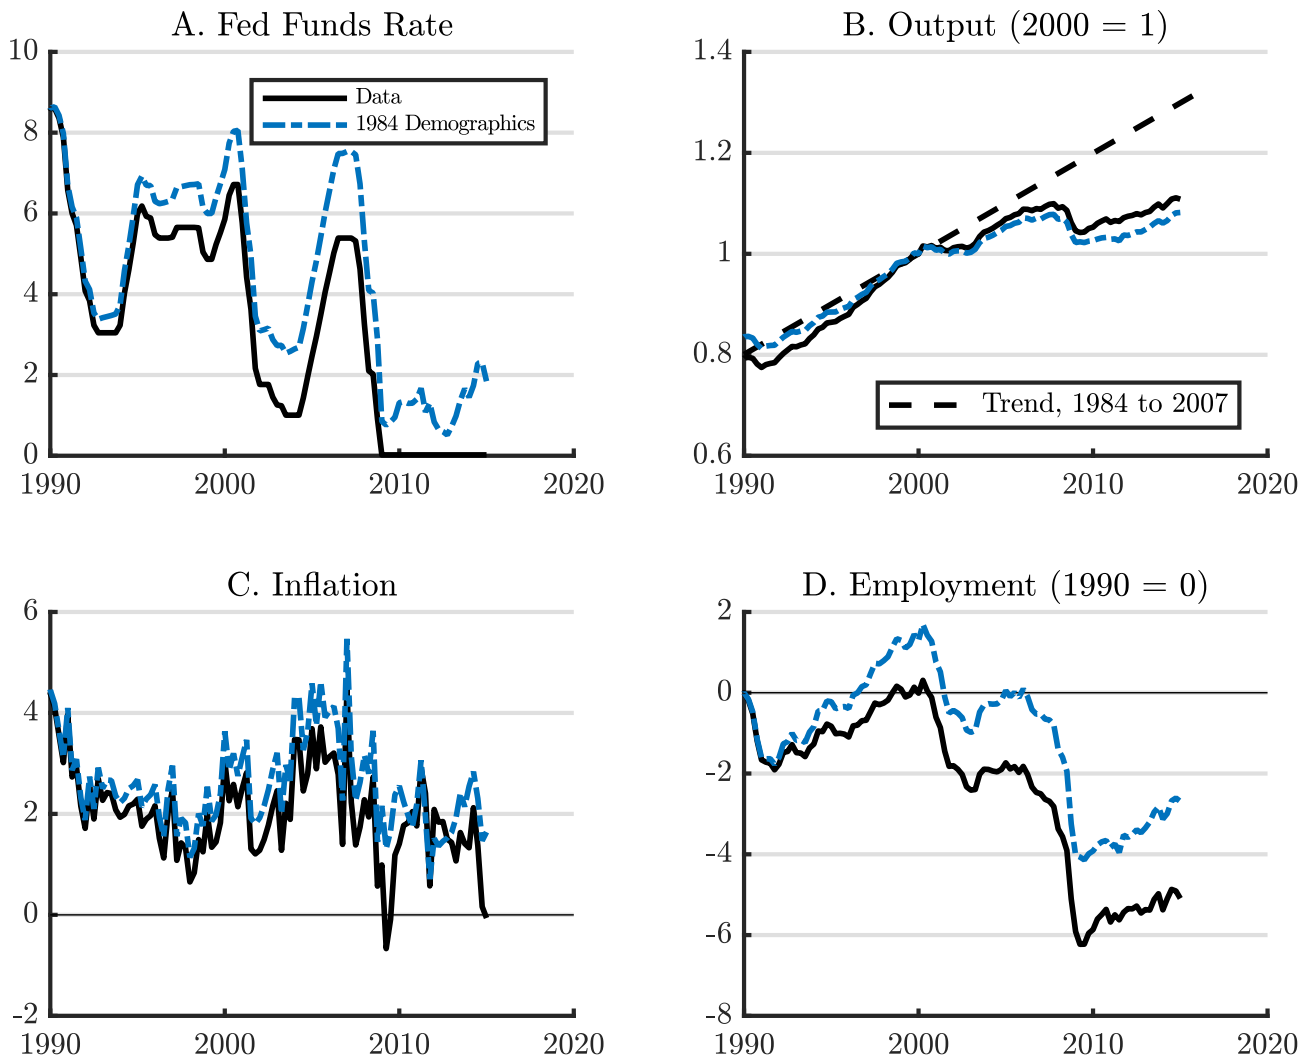


Figure 7: Path of Variables, No Forward Guidance

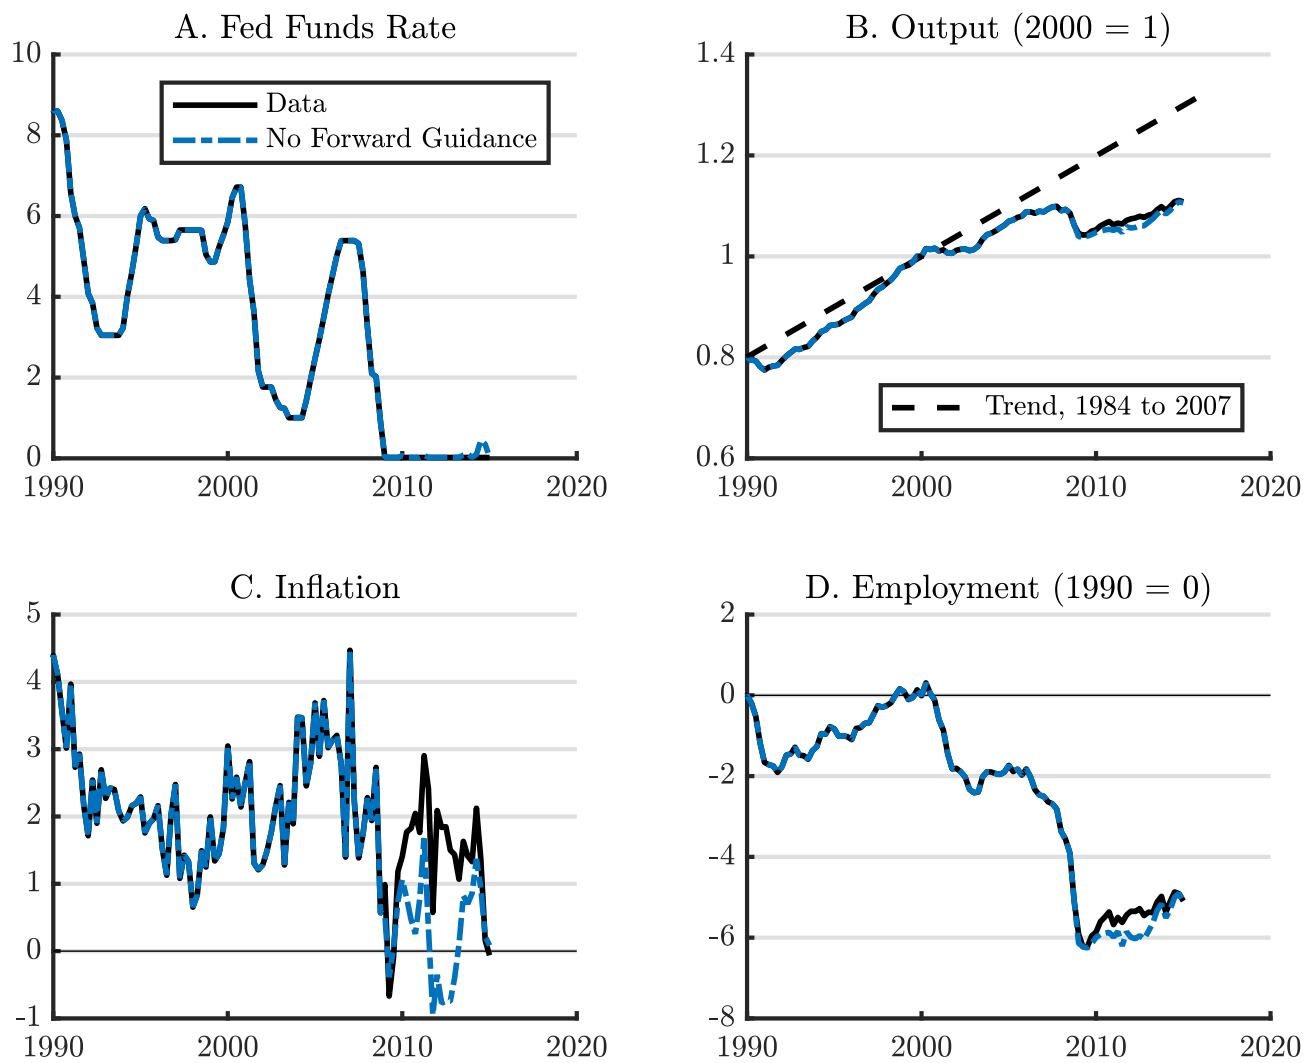


Figure 8: Output Relative to Trend




Figure 9: Output Relative to Trend, Contribution of Shocks

